

TMDb-Movies Dataset Investigation

By: Jeremiah Tindana

Table of Contents

- [Introduction](#)
- [Data Wrangling](#)
- [Exploratory Data Analysis](#)
- [Conclusions](#)

Introduction

Dataset Description

This project details the cleaning, exploration and deduction of a data set from **The Movie Database(Tmdb)**. The data set holds the data concerning 10,000 different movies, along with their user ratings and revenue generated from said movies. The data set is characterized by 21 columns, each holding data on varying characteristics of each movie. The movie title, starring cast, movie director, revenue, runtime, release date and release year are captured.

Per the data set, each movie has its unique ID captured in the **'id'** columns, and an ID from imdb captured in the **'imdb_id'** column.

The **'popularity'** column describes the popularity rating of the movies.

The **'budget'** column captures values equivalent to the monetary value of the resources spent in the creation of the movies. This feature could be analysed to ascertain a relationship between the success of a movie and the movie budget.

The **'tagline'**, **'keywords'** and **'overview'** columns contain brief descriptions of the movie. The tagline captures a single catch phrase that best describes the movie, while the overview contains a paragraph of summary of the movie. **'keywords'** hold a list of words that describe dominant features in the movies. Keywords can be used to describe the genre of each movie. For example, a movie about a monster ravaging a city could have a keyword, 'monster'.

'runtime' holds numeric values equivalent to the length of each movie in minutes. This feature could play an interesting role in the success of a movie. Do movies with shorter runtimes have higher ratings?

'genres' holds a set of genres that the movie falls in. This column is similar in structure to the **'keywords'** column.

'production_companies' holds details on the names of the production companies that cooperated to produce the movie.

'release_date' and **'release_year'** hold the dates of release and year of release of each movie.

Movie patronizers were asked to rate the movies. **'vote_count'** holds the number of patronizers who cast votes. They were asked to rate the movies with scores from 1 to 10, where 10 is the highest score, and 1 is the least score. **'vote_average'** holds the mean rating of each movie.

The final two columns **'budget_adj'** and **'revenue_adj'** show the budget and revenue of the associated movie in terms of 2010 dollars, accounting for inflation over time.

Question(s) for Analysis

The data set in this case study has the prospects to answer a plethora of questions. In this investigation however, we will be limiting ourselves to 3 questions built upon 2 dependent (**popularity** and **revenue**) variables and 3 independent variables (**budget, runtime, and genre**). **Questions:**

1. What is the relationship between the budget of a particular movie and its popularity?
2. How does the popularity of a movie affect its revenue?
3. Have more movies been made over the years? Do older movies have higher equivalent revenue than newer ones?

```
In [1]: # To begin, we import are necessary functions, i.e., pandas, numpy, and matplotlib.  
# We include the matplotlib inline to ensure that our graphs appear in our notebook workspace
```

```
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
%matplotlib inline
```

```
In [2]: # Upgrade pandas to use dataframe.explode() function.  
!pip install --upgrade pandas==0.25.0
```

```
Requirement already up-to-date: pandas==0.25.0 in /opt/conda/lib/python3.6/site-packages (0.25.0)  
Requirement already satisfied, skipping upgrade: numpy>=1.13.3 in /opt/conda/lib/python3.6/site-packages (from pandas==0.25.0) (1.19.5)  
Requirement already satisfied, skipping upgrade: python-dateutil>=2.6.1 in /opt/conda/lib/python3.6/site-packages (from pandas==0.25.0) (2.6.1)  
Requirement already satisfied, skipping upgrade: pytz>=2017.2 in /opt/conda/lib/python3.6/site-packages (from pandas==0.25.0) (2017.3)  
Requirement already satisfied, skipping upgrade: six>=1.5 in /opt/conda/lib/python3.6/site-packages (from python-dateutil>=2.6.1->pandas==0.25.0) (1.11.0)
```

Data Wrangling

In this section, we will take a look at our data and prepare it for analysis. We will scan through the data, carry out some data cleaning where necessary, and trim our data to suit the objectives of this work.

```
In [3]: ## Here, we load our data into this notebook and perform operations to inspect the data.
# We shall acheive this by printing out a few lines of the data and performing a few
# exploratory operations.
df = pd.read_csv("tmdb-movies.csv")
```

```
In [4]: df.head()
```

Out[4]:	id	imdb_id	popularity	budget	revenue	original_title	cast	homepage	director	tagline
0	135397	tt0369610	32.985763	150000000	1513528810	Jurassic World	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	http://www.jurassicworld.com/	Colin Trevorrow	The park open
1	76341	tt1392190	28.419936	150000000	378436354	Mad Max: Fury Road	Tom Hardy Charlize Theron Hugh Keays-Byrne Nic...	http://www.madmaxmovie.com/	George Miller	What Love De
2	262500	tt2908446	13.112507	110000000	295238201	Insurgent	Shailene Woodley Theo James Kate Winslet Ansel...	http://www.thedivergentseries.movie/#insurgent	Robert Schwentke	O Choi C. Destr Y
3	140607	tt2488496	11.173104	200000000	2068178225	Star Wars: The Force Awakens	Harrison Ford Mark Hamill Carrie Fisher Adam D...	http://www.starwars.com/films/star-wars-episod...	J.J. Abrams	Eve generati has sto
4	168259	tt2820852	9.335014	190000000	1506249360	Furious 7	Vin Diesel Paul Walker Jason Statham Michelle ...	http://www.furious7.com/	James Wan	Vengean Hits Hor
5 rows x 21 columns										

```
In [5]: # Let us inspect the data set
df.shape
```

Out[5]:(10866, 21)

With the **df.shape** command, we see the shape of the data frame. The tuple generated from this command shows us the number of rows and columns of our data set. This data set has 10866 rows and 21 columns, indicating that there are, or should be 10866 entries.

```
In [6]: # Let us take a dive into certain statistics about our data
df.describe()
```

Out[6]:	id	popularity	budget	revenue	runtime	vote_count	vote_average	release_year	budget_adj	revenue_adj
count	10866.000000	10866.000000	1.086600e+04	1.086600e+04	10866.000000	10866.000000	10866.000000	10866.000000	1.086600e+04	1.086600e+04
mean	66064.177434	0.646441	1.462570e+07	3.982332e+07	102.070863	217.389748	5.974922	2001.322658	1.755104e+07	5.136436e+07
std	92130.136561	1.000185	3.091321e+07	1.170035e+08	31.381405	575.619058	0.935142	12.812941	3.430616e+07	1.446325e+08
min	5.000000	0.000065	0.000000e+00	0.000000e+00	0.000000	10.000000	1.500000	1960.000000	0.000000e+00	0.000000e+00
25%	10596.250000	0.207583	0.000000e+00	0.000000e+00	90.000000	17.000000	5.400000	1995.000000	0.000000e+00	0.000000e+00
50%	20669.000000	0.383856	0.000000e+00	0.000000e+00	99.000000	38.000000	6.000000	2006.000000	0.000000e+00	0.000000e+00
75%	75610.000000	0.713817	1.500000e+07	2.400000e+07	111.000000	145.750000	6.600000	2011.000000	2.085325e+07	3.369710e+07
max	417859.000000	32.985763	4.250000e+08	2.781506e+09	900.000000	9767.000000	9.200000	2015.000000	4.250000e+08	2.827124e+09

df.describe() presents some statistics about our data. We see our mean, minimum and maximum values. **Popularity:** About 71%, 38%, and 21% of our entries are in the 75th, 50th, and 25th popularity percentile of our data. **Budget:** The maximum budget amount for the movies in our data set is 425 Million, and the minimum budget is 0. The average budget is 146.257 Million. **Revenue:** The maximum revenue generated from any movie in our data set is about 2.78 Billion while the minimum is 0. The average revenue is about 40 Million. **Runtime:** The average runtime for the movies in our data set is about 102 minutes. The longest and shortest movies have 900 minutes and 0 minutes respectively.

```
In [7]: # Let us check for missing entries in our data frame
df.info()
```

```
>class 'pandas.core.frame.DataFrame'>
RangeIndex: 10866 entries, 0 to 10865
Data columns (total 21 columns):
```

```
id                10866 non-null int64
imdb_id           10856 non-null object
popularity        10866 non-null float64
budget            10866 non-null int64
revenue           10866 non-null int64
original_title    10866 non-null object
cast              10790 non-null object
homepage          2936 non-null object
director          10822 non-null object
tagline           8042 non-null object
keywords          9373 non-null object
overview          10862 non-null object
runtime           10866 non-null int64
genres            10843 non-null object
production_companies 9836 non-null object
release_date      10866 non-null object
vote_count        10866 non-null int64
vote_average      10866 non-null float64
release_year      10866 non-null int64
budget_adj        10866 non-null float64
revenue_adj       10866 non-null float64
dtypes: float64(4), int64(6), object(11)
memory usage: 1.7+ MB
```

We see that from the results that certain features have missing data. **imdb_id, cast, homepage, director, tagline, keywords, overview, genres, and production companies**** have missing entries.

Data Cleaning

We can see from the printout from **df.info()** that all data types appear to be in order, except for the **'release_date'** which should preferably be in datetime form. However, the **'release_date'** column in this data is not a key feature as per the objectives of this exercise. For this reason only, we can either ignore or drop this column. All other columns that will play no important role in our analysis can be dropped as well to trim the data, leaving only the necessary features. After dropping the unnecessary columns, the next step will be to handle the null data.

```
In [8]: # Let's take a look at our data again to trim the rows we do not need.
df.head(0)
```

```
Out[8]: id  imdb_id  popularity  budget  revenue  original_title  cast  homepage  director  tagline  ...  overview  runtime  genres  production_companies  rele

0 rows × 21 columns
```



```
In [9]: # We can delete the cast, homepage, director, tagline, overview, production companies and release date columns
# We achieve this with pandas' drop() function.
df.drop(['cast', 'homepage', 'director', 'tagline', 'overview', 'production_companies', 'release_date'], axis=1, inplace=True)
```

```
In [10]: # Let's confirm the changes
df.head()
```

```
Out[10]:
```

	id	imdb_id	popularity	budget	revenue	original_title	keywords	runtime	genres	vote_count
0	135397	tt0369610	32.985763	150000000	1513528810	Jurassic World	monster dna tyrannosaurus rex velociraptor island	124	Action Adventure Science Fiction Thriller	556
1	76341	tt1392190	28.419936	150000000	378436354	Mad Max: Fury Road	future chase post-apocalyptic dystopia australia	120	Action Adventure Science Fiction Thriller	618
2	262500	tt2908446	13.112507	110000000	295238201	Insurgent	novel revolution dystopia sequel dyst...	119	Adventure Science Fiction Thriller	248
3	140607	tt2488496	11.173104	200000000	2068178225	Star Wars: The Force Awakens	android spaceship jedi space opera 3d	136	Action Adventure Science Fiction Fantasy	529
4	168259	tt2820852	9.335014	190000000	1506249360	Furious 7	car race speed revenge suspense car	137	Action Crime Thriller	294



```
In [11]: # Oops! Let's drop the keyword column too. Also, we will not be using the 'budget_adj' column for our analysis
df.drop(['keywords', 'budget_adj'], axis=1, inplace=True)
```

```
In [12]: # confirm changes. We can print only the headers
df.head(0)
```

```
Out[12]: id  imdb_id  popularity  budget  revenue  original_title  runtime  genres  vote_count  vote_average  release_year  revenue_adj
```

```
In [13]: # Take a look at the new shape of the dataframe
df.shape
```

```
Out[13]: (10866, 12)
```

Now that our data frame is neatly trimmed, we have our number of columns reduced from 21 to 11. Next, let's handle null or missing data.

```
In [14]: # Take a look at the rows with missing data
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10866 entries, 0 to 10865
Data columns (total 12 columns):
id          10866 non-null int64
imdb_id     10856 non-null object
popularity  10866 non-null float64
budget      10866 non-null int64
revenue     10866 non-null int64
original_title 10866 non-null object
runtime     10866 non-null int64
genres      10843 non-null object
vote_count  10866 non-null int64
vote_average 10866 non-null float64
release_year 10866 non-null int64
revenue_adj 10866 non-null float64
dtypes: float64(3), int64(6), object(3)
memory usage: 1018.8+ KB

Only 2 features have missing values, i.e., 'imdb_id' and 'genres'. Since the imdb_id values are unique values assigned to each movie, replacing null values with the mean will not be of any benefit. The plausible option here would be to drop those rows. We do same for the genres column.
```

```
In [15]: # Drop null rows using pandas' dropna. Confirm for changes
df.dropna(inplace=True)
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10835 entries, 0 to 10865
Data columns (total 12 columns):
id          10835 non-null int64
imdb_id     10835 non-null object
popularity  10835 non-null float64
budget      10835 non-null int64
revenue     10835 non-null int64
original_title 10835 non-null object
runtime     10835 non-null int64
genres      10835 non-null object
vote_count  10835 non-null int64
vote_average 10835 non-null float64
release_year 10835 non-null int64
revenue_adj 10835 non-null float64
dtypes: float64(3), int64(6), object(3)
memory usage: 1.1+ MB
```

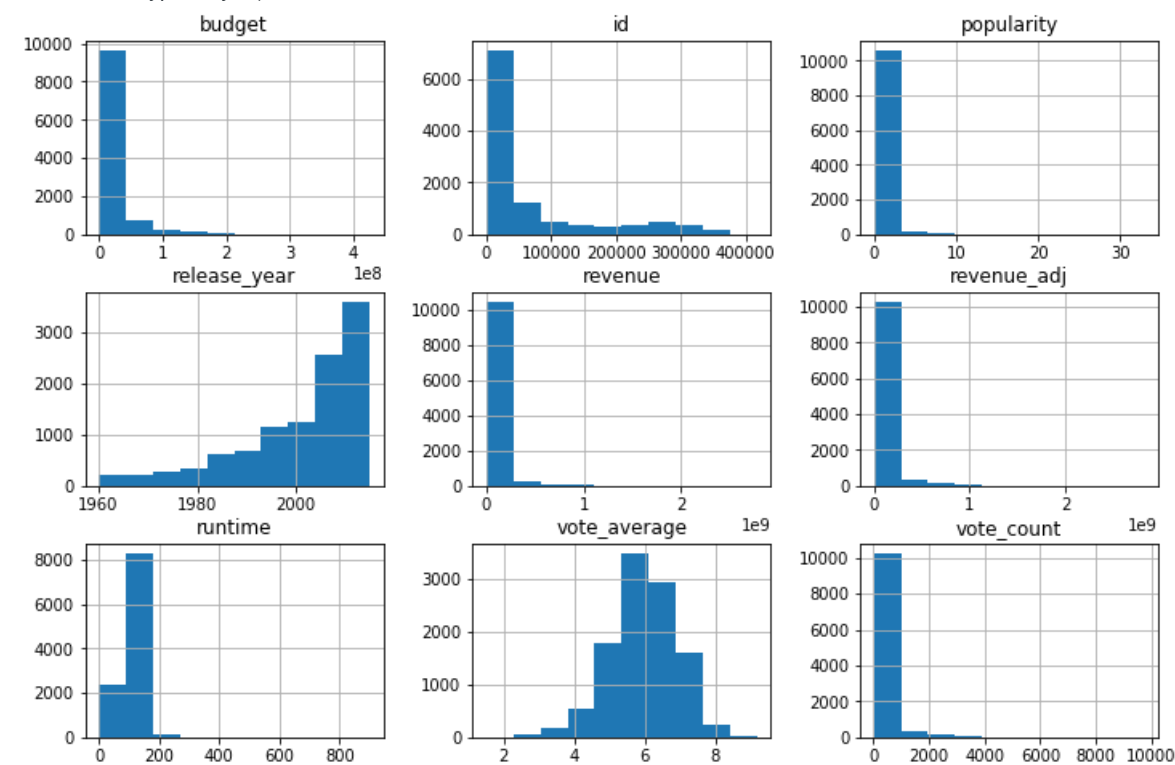
```
In [16]: df.head()
```

Out[16]:	id	imdb_id	popularity	budget	revenue	original_title	runtime	genres	vote_count	vote_average	release_year	rev
0	135397	tt0369610	32.985763	150000000	1513528810	Jurassic World	124	Action Adventure Science Fiction Thriller	5562	6.5	2015	1.39
1	76341	tt1392190	28.419936	150000000	378436354	Mad Max: Fury Road	120	Action Adventure Science Fiction Thriller	6185	7.1	2015	3.48
2	262500	tt2908446	13.112507	110000000	295238201	Insurgent	119	Adventure Science Fiction Thriller	2480	6.3	2015	2.71
3	140607	tt2488496	11.173104	200000000	2068178225	Star Wars: The Force Awakens	136	Action Adventure Science Fiction Fantasy	5292	7.5	2015	1.90
4	168259	tt2820852	9.335014	190000000	1506249360	Furious 7	137	Action Crime Thriller	2947	7.3	2015	1.38

Our data seems to be in pretty good shape. Now, we can delve into our Exploratory Data Analysis. But before that, let's plot some graphical representations of our data just to have an overview of the current state.

```
In [17]: # We'll plot histograms to study our distributions
df.hist(figsize =(12,8))
```

```
Out[17]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x7f10d892d5f8>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7f10d891de48>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7f10d88d8dd8>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x7f10d88c2390>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7f10d884f0b8>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7f10d884f0f0>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x7f10d87bf128>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7f10d8777208>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7f10d87a1908>]],
dtype=object)
```



Interpreting the Graphs

The graphs give us some comfort that our analysis has so far been on the right path. Our graphs make a lot of sense, and we see no severe deviations or abnormalities. Most of our graphs are skewed. The greater part of the movie budgets are less than 1 billion. The popularity of most of the movies is between 0 and 10. The graph for the release year is skewed to the left, indicating that most of the movies were released after the year 2000, and a number of them between 1960 and 2000. The graph of revenues is very skewed to the right. Almost all our movies have revenues less than 500 Million. Only a few movies have runtimes above 200 minutes.

Exploratory Data Analysis


1. What is the relationship between the budget of a particular movie and its popularity?

Dependent Variable - Popularity

First of all let us plot a scatter graph of **budget** and **popularity** to check for correlation. Since we will be performing this task a number of times in our analysis, we need to make this task less cumbersome. We will achieve this by creating a function.

```
In [18]: # Create a function to plot scatter graphs
def scatter_graph(arg):
    df.plot(x='popularity', y=arg, kind='scatter')

scatter_graph('budget')
```



Our scatter shows a level of positive correlation between budget and popularity. We can explore this relationship even further.

```
In [19]: # Let's take a look at our dependent variable
```

```
df.popularity.max(), df.popularity.median(), df.popularity.min())
```

```
Out[19]:(32.985763, 0.38461799999999996, 6.500000000000001e-05)
```

Our maximum, median and minimum popularity are as shown above. Let's look at our median popularity.

```
In [20]: median = df.popularity.median()
         median
```

```
Out[20]:0.38461799999999996
```

Let's see how many movies had popularity above the median popularity

```
In [21]: popular = df.query('popularity>{}'.format(median))
         unpopular = df.query('popularity<{}'.format(median))
```

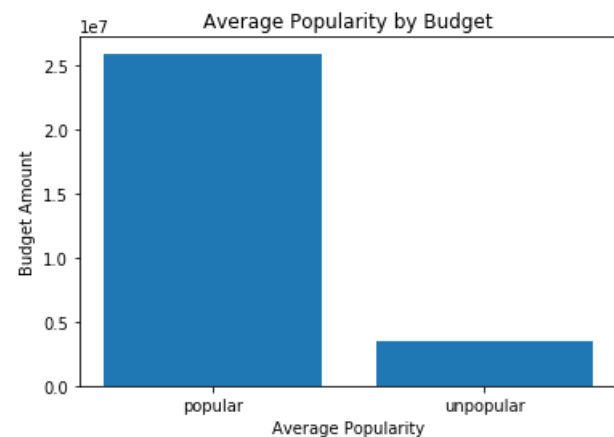
```
In [22]: # Let's find the mean for each group in relation to budget
         mean_popular = popular['budget'].mean()
         mean_unpopular = unpopular['budget'].mean()
         mean_popular, mean_unpopular
```

```
Out[22]:(25892745.875392284, 3443115.780321211)
```

Now that we have our means, let us create a bar chart to visualize the relationships in our data.

```
In [23]: # Creating a bar chart
         locations = [1,2]
         heights = [mean_popular, mean_unpopular]
         labels = ['popular','unpopular']
         plt.bar(locations,heights,tick_label=labels)
         plt.title('Average Popularity by Budget')
         plt.xlabel('Average Popularity')
         plt.ylabel('Budget Amount')
```

```
Out[23]:Text(0,0.5,'Budget Amount')
```




The bar chart provides more conclusive and substantial results. The most popular movies appear to have higher budget amounts than the unpopular movies. The budget of the movies whose mean popularity are above the median are greater than those whose popularity are below the median.

2. How does the popularity of a movie affect its revenue?

Like before, we shall begin by checking for correlation using a scatter plot

```
In [24]: # Creating a scatter plot of revenue and popularity
         scatter_graph('revenue')
```



It would make sense to assume that the most popular movies would have higher revenues. However, let's not be quick to assume. The scatter graph shows some positive correlation. We can go further to get more assertive results.

```
In [25]: # Take a look at some descriptive statistics of our revenue data
         df.revenue.describe()
```

```
Out[25]:count    1.083500e+04
        mean     3.993726e+07
        std      1.171513e+08
        min      0.000000e+00
        25%      0.000000e+00
        50%      0.000000e+00
        75%      2.417286e+07
        max      2.781506e+09
        Name: revenue, dtype: float64
```

The mean revenue is 39937260 (about 40 Million)

```
In [26]: mean_rev = df.revenue.mean()
```

We can go ahead to divide the revenue set into different levels. We can create a new column called **revenue_level** and place in there the level of revenue generated, either **'high'** or **'low'**. The low and high categories will be created with regards to the mean.

```
In [27]: # We use a for loop to create our new column in the data frame
```

```
revenue_level = []
for value in df['revenue']:
    if value >= mean_rev:
        revenue_level.append('high')
    else:
        revenue_level.append('low')
```

```
df['revenue_level'] = revenue_level
```

```
# We need to check if our changes have taken effect
df.revenue_level
```

```
Out[27]:0    high
        1    high
        2    high
        3    high
        4    high
        ...
        10861 low
        10862 low
        10863 low
        10864 low
        10865 low
        Name: revenue_level, Length: 10835, dtype: object
```


Since our changes have taken effect, we can proceed.

```
In [28]: # Let's proceed to fish out our two groups
```

```
high_rev = df[df['revenue'] >= mean_rev]
low_rev = df[df['revenue'] < mean_rev]
```

```
In [29]: high_rev['popularity'].hist(alpha=0.5,bins=20,label='Popular')
        low_rev['popularity'].hist(alpha=0.5,bins=20,label='Unpopular')
        plt.legend()
```

```
Out[29]:<matplotlib.legend.Legend at 0x7f10d7700a58>
```



Look at that! Our data does not agree with the initial assumption! Looking at our double plot, we can see that on the average, per our data set, the unpopular movies seem to have a higher average revenue than the popular ones. That is surprising isn't it? Well it can be explained. In the two new data sets we created, **high_rev** and **low_rev**, we see that our criteria for splitting the data wasn't the most appropriate. That is because our data is skewed, and most of the entries in the data set are on the low side. This would mean that even if the popular ones have higher revenues, the overwhelming number of unpopular ones would cause inaccurate results. A solution to this would be to create more than 2 groups, to ensure that the data is split somewhat evenly.


3. Have more movies been made over the years? Do older movies have higher equivalent revenue than newer ones?

With this question, we want to compare the years in our movie data by checking their revenues. We must take note however that the value of money changes from year to year. Luckily we have the **revenue_adj** column which is the monetary equivalent in 2010 for all the years.

```
In [30]: # Let's view our data again
        df.head(2)
```

Out[30]:

	id	imdb_id	popularity	budget	revenue	original_title	runtime	genres	vote_count	vote_average	release_year	rev
0	135397	tt0369610	32.985763	150000000	1513528810	Jurassic World	124	Action Adventure Science Fiction Thriller	5562	6.5	2015	1.39
1	76341	tt1392190	28.419936	150000000	378436354	Mad Max: Fury Road	120	Action Adventure Science Fiction Thriller	6185	7.1	2015	3.48



```
In [31]: # Let's explore our release_year column
df.release_year.describe()
```

```
Out[31]:count    10835.000000
         mean      2001.308999
         std       12.815519
         min       1960.000000
         25%       1995.000000
         50%       2006.000000
         75%       2011.000000
         max       2015.000000
         Name: release_year, dtype: float64
```

So the least year of release captured in our data is 1960, and the latest year is 2015. Let's look at the unique years in our data set.

```
In [32]: # we will use the unique function
df['release_year'].nunique()
```

```
Out[32]:56
To break the data down into years, we can look at the amount of data in each year. We will use the value_counts() function
```

```
In [33]: df['release_year'].value_counts()
```



```

Out[33]:2014 699
        2013 655
        2015 626
        2012 583
        2011 540
        2009 529
        2008 495
        2010 487
        2007 436
        2006 408
        2005 363
        2004 307
        2003 281
        2002 266
        2001 241
        2000 226
        1999 224
        1998 210
        1996 203
        1997 192
        1994 184
        1993 178
        1995 174
        1988 145
        1989 136
        1992 133
        1991 133
        1990 132
        1987 125
        1986 121
        1985 109
        1984 105
        1981 82
        1982 81
        1983 80
        1980 78
        1978 65
        1979 57
        1977 57
        1971 55
        1973 55
        1976 47
        1974 46
        1966 46
        1975 44
        1964 42
        1970 40
        1967 40
        1972 40
        1968 39
        1965 35
        1963 34
        1962 32
        1960 32
        1969 31
        1961 31

```


Name: release_year, dtype: int64

Here, we can see that the movies across the years from 1961 forward keep increasing. Let's plot a histogram to visualize this trend.

```
In [34]: plt.hist(df['release_year'])
```

```

Out[34]:(array([ 206., 196., 287., 339., 621., 679., 1141., 1238., 2538.,
                3590.]),
         array([1960. , 1965.5, 1971. , 1976.5, 1982. , 1987.5, 1993. , 1998.5,
                2004. , 2009.5, 2015. ]),
         <a list of 10 Patch objects>)
```



Skewed data! Our graph of release years is skewed to the left. This proves that as the years went by, more and more movies were produced. The movie industry has grown largely over the decades. But by how much?

```
In [35]: # Find the percentage increase between 1961 and 2014. Select data from release year value counts
mov_1961 = 31
mov_2014 = 699
perc_inc = (mov_2014 / mov_1961) * 100
perc_inc
```


Out[35]:2254.838709677419

The movie industry has grown by over 2000%!

Now let's take a look at the relationship between revenue_adj and the movies across the years

```
In [36]: plt.bar(df['release_year'],df['revenue_adj'])
plt.title('2010-Equivalent Revenue for movies')
plt.xlabel('Revenue')
plt.ylabel('Year')
```

Out[36]:Text(0,0.5,'Year')



No conclusive results can be drawn from the graph. Converting all revenue values from all years to equivalent values of money, each year had its revenue independent. The revenue generated from movies depends on some other factors, but not the year.

Conclusions

1. What is the relationship between the budget of a particular movie and its popularity?

Movies with higher budgets tend to end up with more popularity than those with low budgets.

2. How does the popularity of a movie affect its revenue?

The group of unpopular movies seem to have a higher average revenue than the popular ones.

3. Have more movies been made over the years? Do older movies have higher equivalent revenue than newer ones?

The movie industry has grown by over 2000%! Older movies did not appear to have more or less equivalent revenue than the new ones.

Limitations

4. In the two new data sets created in the second analysis, high_rev and low_rev, we see that our criteria for splitting the data wasn't the most appropriate. That is because our data is skewed, and most of the entries in the data set are on the low side. This would mean that even if the popular ones have higher revenues, the overwhelming number of unpopular ones would cause inaccurate results. A solution to this would be to create more than 2 groups, to ensure that the data is split almost evenly.
5. The same can be said about the popular and unpopular data sets in the first analysis.

```
In [38]: from subprocess import call
call(['python', '-m', 'nbconvert', 'Investigate_a_Dataset.ipynb'])
```

Out[38]:0

In []: